

# Statistics Review

EC320, Set 02

Andrew Dickinson

Spring 2024



# Admin

# Admin

## 1. and RStudio Install

## 2. Lab

- Lab module on Canvas homepage
- First recording available now

## 1. Koans

- First Koans due next Friday
- Get started now

Stopping point after lecture 02

## 1. PS01

- First problem set will be assigned next week
- Due next Tuesday (04/16)

## 1. Textbook

# Motivation

The focus of our course is **regression analysis**—part of the fundamental toolkit for learning from data.

The **underlying theory** is critical to grasp the mechanics and pitfalls

- Make us better practitioners and savvier consumers of science.

**Today:** Review the essential concepts from Math 243

# Warning.

The following review is a lot packed in very briefly though you *should* have learned much of it before. But that being said, it will be overwhelming for most.

# Notation

# Notation

Data on a variable  $X$  are a sequence of  $n$  observations, indexed by  $i$ :

$$\{x_i : 1, \dots, n\}.$$

**Ex.**  $n = 5$

$i$	$x_i$
1	8
2	9
3	4
4	7
5	2

- $i$  indicates the row number.
- $n$  is the number of rows.
- $x_i$  is the value of  $X$  for row  $i$ .



# Summation

The **summation operator** adds a sequence of numbers over an index:

$$\sum_{i=1}^n x_i \equiv x_1 + x_2 + \cdots + x_n.$$

The sum of  $x_i$  from 1 to  $n$ .

$i$	$x_i$
1	7
2	4
3	10
4	3

$$\sum_{i=1}^4 x_i = 7 + 4 + 10 + 3 = 23$$

$$\frac{1}{n} \sum_{i=1}^n x_i \rightarrow \frac{1}{4} \sum_{i=1}^4 x_i = 6$$

# Summation

The **summation operator** adds a sequence of numbers over an index:

$$\sum_{i=1}^n x_i \equiv x_1 + x_2 + \cdots + x_n.$$

The sum of  $x_i$  from 1 to  $n$ .

$i$	$c$
1	2
2	2
3	2
4	2

$$\sum_{i=1}^4 x_i = 7 + 4 + 10 + 3 = 23$$

$$\text{sample average} \left\{ \frac{1}{n} \sum_{i=1}^n x_i \rightarrow \frac{1}{4} \sum_{i=1}^4 x_i = 6 \right.$$

EC320, Set 02 | Statistics Review

# Summation: Rule 01

For any constant  $c$ ,

$$\sum_{i=1}^n c = nc.$$

$i$	$c$
1	2
2	2
3	2
4	2

$$\begin{aligned}\sum_{i=1}^4 2 &= 4 \times 2 \\ &= 8\end{aligned}$$

# Summation: Rule 02

For any constant  $c$ ,

$$\sum_{i=1}^n cx_i = c \sum_{i=1}^n x_i.$$

$i$	$x_i$
1	8
2	9
3	4
4	7
5	2

$$\begin{aligned} \sum_{i=1}^3 2x_i &= 2 \times 7 + 2 \times 4 + 2 \times 10 \\ &= 14 + 8 + 20 = 42 \end{aligned}$$

$$2 \sum_{i=1}^3 x_i = 2(7 + 4 + 10) = 42$$

# Summation: Rule 03

If  $\{(x_i, y_i) : 1, \dots, n\}$  is a set of  $n$  pairs, and  $a$  and  $b$  are constants, then

$$\sum_{i=1}^n (ax_i + by_i) = a \sum_{i=1}^n x_i + b \sum_{i=1}^n y_i$$

$i$	$a$	$x_i$	$b$	$y_i$
1	2	7	1	4
2	2	4	1	2

$$\sum_{i=1}^2 (2x_i + y_i) = 18 + 10 = 28 \quad (1)$$

$$2 \sum_{i=1}^2 x_i + \sum_{i=1}^2 y_i = 2 \times 11 + 6 = 28 \quad (2)$$

# Summation: Caution 01

The **sum of the ratios** is not the **ratio of the sums**:

$$\sum_{i=1}^n x_i / y_i \neq \left( \sum_{i=1}^n x_i \right) / \left( \sum_{i=1}^n y_i \right)$$

**Ex.**

If  $n = 2$ , then  $\frac{x_1}{y_1} + \frac{x_2}{y_2} \neq \frac{x_1 + x_2}{y_1 + y_2}$

# Summation: Caution 02

The **sum of squares** is not the **square of the sums**:

$$\sum_{i=1}^n x_i^2 \neq \left( \sum_{i=1}^n x_i \right)^2$$

**Ex.**

If  $n = 2$ , then  $x_1^2 + x_2^2 \neq (x_1 + x_2)^2 = x_1^2 + 2x_1x_2 + x_2^2$ .

# Cartesian coordinate system

Cartesian plane: 2-D plane defined by two perpendicular number lines:

- x-axis (*horizontal*)
- y-axis (*vertical*)

Using these axes, any point in the plane is described using an ordered pair of numbers  $(x, y)$



# Cartesian coordinate system

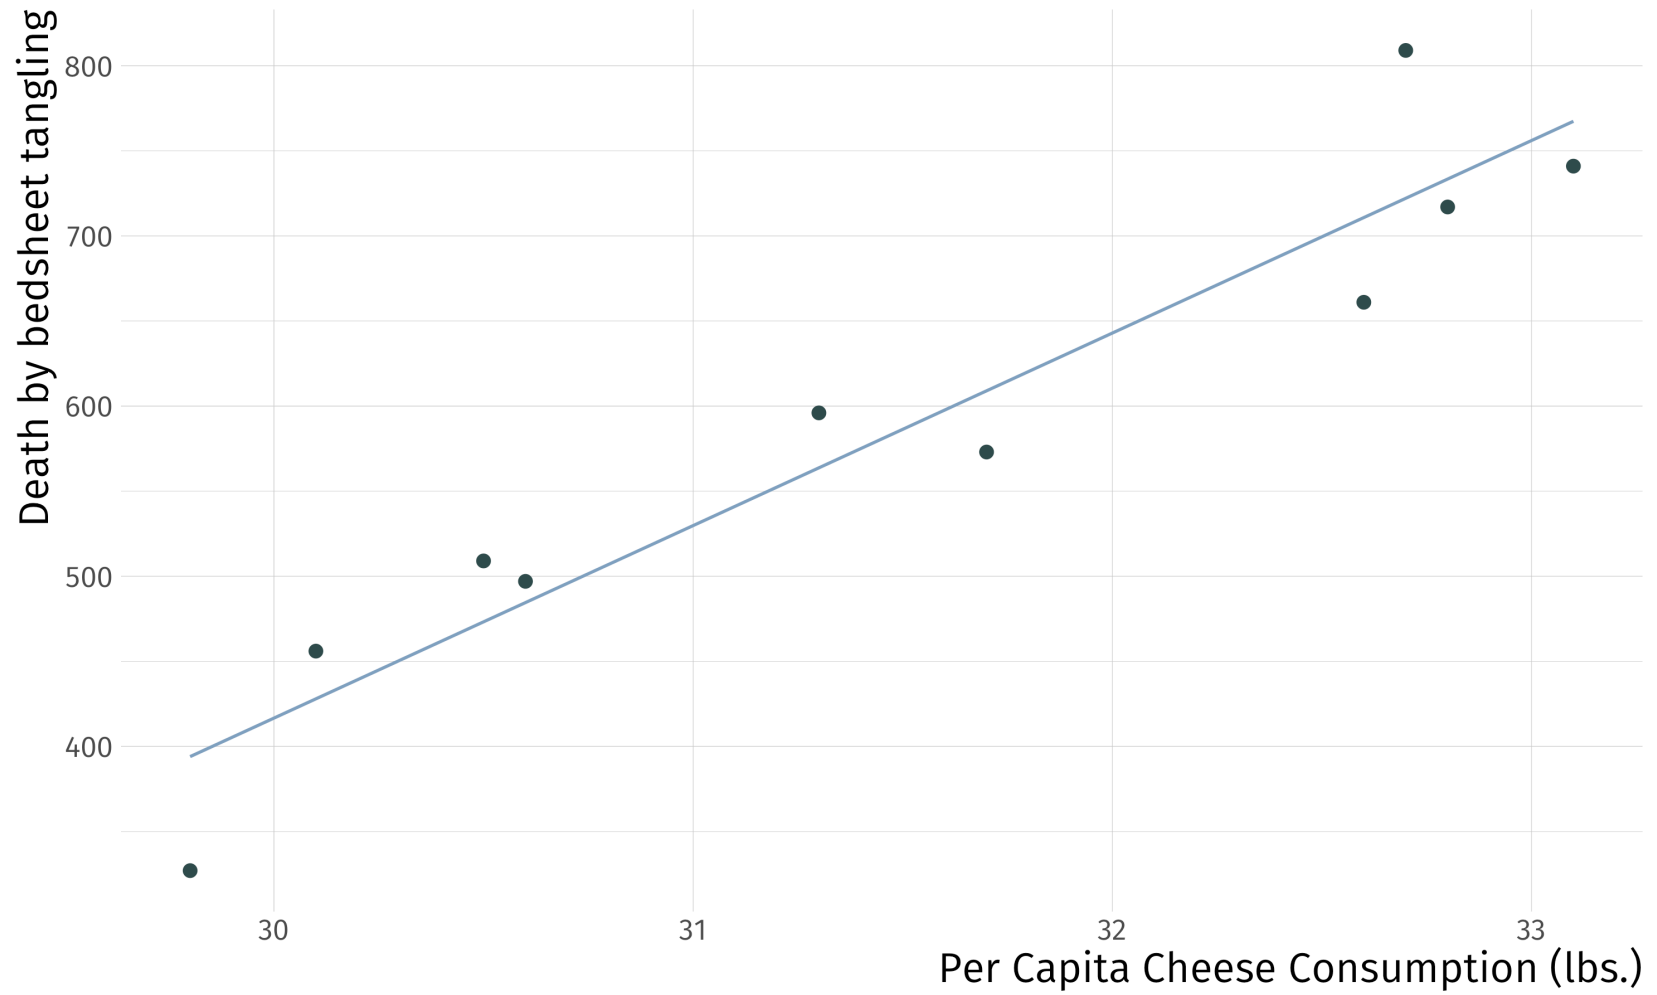
A particular line on this plane takes the form

$$y = a + bx$$

where *a* is known as the intercept and *b* is the slope.

Any incremental unit increase in *x* results in *y* increasing by *b*.

**Ex.**



# Basic probability

# Essential definitions

## Experiment:

Any procedure that is *infinitely repeatable* and has a *well-defined set of outcomes*.

**Ex.** Flip a coin 10 times and record the number of heads.

## Random Variable:

A variable with *numerical values determined by an experiment or a random phenomenon*.

- Describes the sample space of an experiment.

# Essential definitions

## Sample Space:

The set of potential outcomes an experiment could generate

**Ex.** The sum of two dice is an integer from 2 to 12.

## Event:

A subset of the sample space or a combination of outcomes.

**Ex.** Rolling a two or a four.

# Random variables

**Notation:** Capital letters for random variables (e.g.,  $X$ ,  $Y$ , or  $Z$ ) and lowercase letters for particular outcomes (e.g.,  $x$ ,  $y$ , or  $z$ ).

## Experiment

Flipping a coin.

## Events:

Heads or tails.

## Random Variable: ( $X$ )

Receive \$1 if heads,  $x_i = 1$ , pay \$1 if tails,  $x_i = -1$

## Sample Space:

$\{-1, 1\}$

# Discrete random variables

A random variable that takes a countable set of values.

## **Bernoulli (*binary*) random variable**

Random variable that takes values of either 1 or 0.

- Characterized by  $P(X = 1)$ , “the probability of success.”
- Probabilities sum to 1:  $P(X = 1) + P(X = 0) = 1$

More generally, if

then

$$P(X = 1) = \theta$$

$$P(X = 0) = 1 - \theta$$

for some  $\theta \in [0, 1]$

# Discrete Random Variables: Probabilities

We describe a discrete random variable by listing its possible values with associated probabilities.

If  $X$  takes on  $k$  possible values  $\{x_1, \dots, x_k\}$ , then the probabilities  $p_1, p_2, \dots, p_k$  are defined by

$$p_j = P(X = x_j), \quad j = 1, 2, \dots, k,$$

where

$$p_j \in [0, 1]$$

and

$$p_1 + p_2 + \dots + p_k = 1.$$



# Discrete Random Variables

## Probability density function (pdf)

The *pdf* of  $X$  summarizes possible outcomes and associated probabilities:

$$f(x_j) = p_j, \quad j = 1, 2, \dots, k.$$

**Ex.** 2020 Presidential election: 538 electoral votes at stake.

- $\{X : 0, 1, \dots, 538\}$  is the number of votes won.
- Unlikely that one will win 0 or 538 votes:  $f(0) \approx 0$  and  $f(538) \approx 0$ .
- Nonzero probability of winning an exact majority:  $f(270) > 0$ .

# Discrete random variables **Ex.**

Basketball player goes to the foul line to shoot two free throws.

- $X$  is the number of shots made (either 0, 1, or 2).
- The pdf of  $X$  is  $f(0) = 0.3$ ,  $f(1) = 0.4$ ,  $f(2) = 0.3$ .<sup>1</sup>

Use the pdf to calculate the probability of the **event** that the player makes *at least one shot*, i.e.,  $P(X \geq 1)$ .

$$P(X \geq 1) = P(X = 1) + P(X = 2) = 0.4 + 0.3 = 0.7$$

<sup>1</sup> Note: the probabilities sum to 1

# Continuous random variables

A random variable that takes any real value with *zero* probability.

*Wait, what?!* The variable takes so many values that we can't count all possibilities, so the probability of any one particular value is zero.

Measurement is discrete (*e.g.*, dollars and cents), but variables with many possible values are best treated as continuous.

- *e.g.*, electoral votes, height, wages, temperature, *etc.*

# Continuous random variables

Probability density functions also describe continuous random variables.

Difference between continuous and discrete **PDFs**

- Interested in the probability of events within a *range* of values.
- *e.g.* What is the probability of more than 1 inch of rain tomorrow?

# Distributions

# Distributions

Function that represents all outcomes of a random variable and the corresponding probabilities.


- Summary that describes the spread of data points in a set
- Essential for making inferences and assumptions from data

**Key Takeaway:** The shape of a distribution provides valuable information

# Uniform distribution

The probability density function of a variable uniformly distributed between 0 and 2 is

$$f(x) = \begin{cases} \frac{1}{2} & \text{if } 0 \leq x \leq 2 \\ 0 & \text{otherwise} \end{cases}$$




# Uniform distribution

By definition, the area under  $f(x)$  is equal to 1.

The **shaded area** illustrates the probability of the event  $1 \leq X \leq 1.5$ .

$$P(1 \leq X \leq 1.5) = (1.5 - 1) \times 0.5 = 0.25$$





# Normal Distribution

The “**bell curve**”

- Symmetric: mean and median occur at the same point (*i.e.*, no skew).
- Low-probability events in tails; high-probability events near center.




# Normal Distribution

The **shaded area** illustrates the probability of the event  $-2 \leq X \leq 2$ .

- “Find area under curve” = use integral calculus (or, in practice, **R**).

$$P(-2 \leq X \leq 2) \approx 0.95$$



# Normal Distribution

Continuous distribution where  $x_i$  takes the value of any real number ( $\mathbb{R}$ )

- Domain spans the entire real line
- Centered on the distribution mean  $\mu$

Rule 1: The probability that the random variable takes a value  $x_i$  is 0 for any  $x_i \in \mathbb{R}$

Rule 2: The probability that the random variable falls between  $[x_i, x_j]$  range, where  $x_i \neq x_j$ , is the area under  $p(x)$  between those two values

The area above represents  $p(x) = 0.95$ . The values  $\{-1.96, 1.96\}$  represent the 95% confidence interval for  $\mu$ .

# Moments

# Moments

Quantitative measures used to describe the shape and characteristics of a probability distribution<sup>1</sup>

Summarize and understand the important features of a distribution

First moment: **Mean**

Second moment: **Variance**

Third moment: Skewness

Fourth moment: Kurtosis

⋮

# Expected Value

Describes the *central tendency* of distribution in a single number.<sup>1</sup>

Density functions describe the entire distribution, but sometimes we just want a summary.

Other summary statistics we may be interested in include

- Median
- Standard deviation
- 25th percentile
- 75th percentile

# Expected Value (discrete)

The expected value of a discrete random variable  $X$  is the weighted average of its  $k$  values  $\{x_1, \dots, x_k\}$  and their associated probabilities:

$$\begin{aligned} E(X) &= x_1P(x_1) + x_2P(x_2) + \dots + x_kP(x_k) \\ &= \sum_{j=1}^k x_jP(x_j). \end{aligned}$$

AKA: **Population mean**

## Expected Value **Ex.**

Rolling a six-sided die once can take values  $\{1, 2, 3, 4, 5, 6\}$ , each with equal probability. *What is the expected value of a roll?*

$$\begin{aligned} E(\text{Roll}) &= 1 \times \frac{1}{6} + 2 \times \frac{1}{6} + 3 \times \frac{1}{6} + 4 \times \frac{1}{6} \\ &\quad + 5 \times \frac{1}{6} + 6 \times \frac{1}{6} = 3.5 \end{aligned}$$


Note: The **EV** can be a number that isn't a possible outcome of ***X***.



# Expected value (continuous)

If  $X$  is a continuous random variable and  $f(x)$  is its probability density function, then the expected value of  $X$  is

$$E(X) = \int_{-\infty}^{\infty} x f(x) dx.$$



# Expected value: Rule 01

For any constant  $c$ ,  $E(c) = c$ . **Ex.**

- $E(5) = 5$ .
- $E(1) = 1$ .
- $E(4700) = 4700$ .

# Expected value: Rule 02

For any constants  $a$  and  $b$ ,  $E(aX + b) = aE(X) + b$ .

**Ex.** Suppose  $X$  is the high temperature in degrees Celsius in Eugene during August. The long-run average is  $E(X) = 28$ . If  $Y$  is the temperature in degrees Fahrenheit, then  $Y = 32 + \frac{9}{5}X$ . What is  $E(Y)$ ?

$$E(Y) = 32 + \frac{9}{5}E(X) = 32 + \frac{9}{5} \times 28 = 82.4$$

# Expected value: Rule 03

If  $\{a_1, a_2, \dots, a_n\}$  are constants and  $\{X_1, X_2, \dots, X_n\}$  are random variables, then

$$E(a_1X_1 + a_2X_2 + \dots + a_nX_n) = a_1E(X_1) + a_2E(X_2) + \dots + a_nE(X_n)$$

In English, **the expected value of the sum = the sum of expected values.**

# Expected value: Rule 03

**The expected value of the sum = the sum of expected values.**

**Ex.** Suppose that a coffee shop sells  $X_1$  small,  $X_2$  medium, and  $X_3$  large caffeinated beverages in a day. The quantities sold are random with expected values  $E(X_1) = 43$ ,  $E(X_2) = 56$ , and  $E(X_3) = 21$ . The prices of small, medium, and large beverages are 1.75, 2.50, and 3.25 dollars. What is expected revenue?

$$\begin{aligned} E(1.75X_1 + 2.50X_2 + 3.35X_3) &= 1.75E(X_1) + 2.50E(X_2) + 3.25E(X_3) \\ &= 1.75(43) + 2.50(56) + 3.25(21) \\ &= 283.5 \end{aligned}$$

# Expected value: Caution

Previously, we found that the expected value of rolling a six-sided die is  $E(\text{Roll}) = 3.5$ .

- If we square this number, we get  $[E(\text{Roll})]^2 = 12.25$ .

Is  $[E(\text{Roll})]^2$  the same as  $E(\text{Roll}^2)$ ?

$$\begin{aligned} E(\text{Roll}^2) &= 1^2 \times \frac{1}{6} + 2^2 \times \frac{1}{6} + 3^2 \times \frac{1}{6} + 4^2 \times \frac{1}{6} \\ &\quad + 5^2 \times \frac{1}{6} + 6^2 \times \frac{1}{6} \\ &\approx 15.167 \neq 12.25. \end{aligned}$$

**No!**

# Expected value: Caution


Except in special cases, **the transformation of an expected value** is not **the expected value of a transformed random variable**.

For some function  $g(\cdot)$ , it is typically the case that

$$g(E(X)) \neq E(g(X)).$$

# Variance

Random variables  $X$  and  $Y$  share the same population mean, but are distributed differently.





# Variance ( $\sigma^2$ )

Tells us how far  $X$  deviates from  $\mu$ , *on average*:

$$\text{Var}(X) \equiv E[(X - \mu)^2] = \sigma_X^2$$

Where:  $\mu = E(X)$ .

*How tightly is a random variable distributed about its mean?*

Describe the distance of  $X$  from its population mean  $\mu$  as the squared difference:  $(X - \mu)^2$ .

- Distributing the terms above yields  $\sigma^2 = E(X^2 - 2X\mu + \mu^2) = E(X^2) - 2\mu^2 + \mu^2 = E(X^2) - \mu^2$ .

# Variance: Rule 01

$\text{Var}(X) = 0 \iff X$  is a constant.

- A random variable that never deviates from its mean has zero variance.

*Wait what? How can a random variable be a constant??* Because a constant fits the technical definition of a random variable<sup>1</sup>. It's just not-so-random

# Variance: Rule 02

For any constants  $a$  and  $b$ ,  $\text{Var}(aX + b) = a^2 \text{Var}(X)$ .

**Ex.** Suppose  $X$  is the high temperature in degrees Celsius in Eugene during August. If  $Y$  is the temperature in degrees Fahrenheit, then  $Y = 32 + \frac{9}{5}X$ . What is  $\text{Var}(Y)$ ?

$$\text{Var}(Y) = \left(\frac{9}{5}\right)^2 \text{Var}(X) = \frac{81}{25} \text{Var}(X)$$

# Standard Deviation ( $\sigma$ )

The positive square root of the variance:

$$\text{sd}(X) = +\sqrt{\text{Var}(X)} = \sigma$$

**Rule 01:** For any constant  $c$ ,  $\text{sd}(c) = 0$ .

**Rule 02:** For any constants  $a$  and  $b$ ,  $\text{sd}(aX + b) = |a| \text{sd}(X)$ .

Note: The same as variance, almost

# Standardizing a random variable

When we're working with a random variable  $X$  with an unfamiliar scale, it is useful to **standardize** it by defining a new variable  $Z$ :

$$Z \equiv \frac{X - \mu}{\sigma}$$

$Z$  has mean **0** and standard deviation **1**. How?

- First, some simple trickery:  $Z = aX + b$ , where  $a \equiv \frac{1}{\sigma}$  and  $b \equiv -\frac{\mu}{\sigma}$ .
- $E(Z) = aE(X) + b = \mu \frac{1}{\sigma} - \frac{\mu}{\sigma} = 0$ .
- $\text{Var}(Z) = a^2 \text{Var}(X) = \frac{1}{\sigma^2} \sigma^2 = 1$ .

# Covariance

For two random variables  $X$  and  $Y$ , the covariance is defined as the expected value (or mean) of the product of their deviations from their individual expected values:

$$\text{Cov}(X, Y) \equiv E[(X - \mu_X)(Y - \mu_Y)] = \sigma_{XY}$$

**Idea:** Characterize the relationship between random variables  $X$  and  $Y$ .

- **Positive correlation:** When  $\sigma_{XY} > 0$ , then  $X$  is above its mean when  $Y$  is above its mean, *on average*.
- **Negative correlation:** When  $\sigma_{XY} < 0$ , then  $X$  is below its mean when  $Y$  is above its mean, *on average*.

# Covariance: Rule 01

## Statistical independence:

If  $X$  and  $Y$  are independent, then  $E(XY) = E(X)E(Y)$ .

- If  $X$  and  $Y$  are independent, then  $\text{Cov}(X, Y) = 0$ .

Caution:

- $\text{Cov}(X, Y) = 0$  **does not imply** that  $X$  and  $Y$  are independent.
- $\text{Cov}(X, Y) = 0$  means that  $X$  and  $Y$  are *uncorrelated*.

# Covariance: Rule 02

For any constants  $a$ ,  $b$ ,  $c$ , and  $d$ ,

$$\text{Cov}(aX + b, cY + d) = ac \text{Cov}(X, Y)$$



# Correlation Coefficient


A problem with covariance is that it is sensitive to units of measurement. The **correlation coefficient** solves this problem by rescaling the covariance:

$$\text{Corr}(X, Y) \equiv \frac{\text{Cov}(X, Y)}{\text{sd}(X) \times \text{sd}(Y)} = \frac{\sigma_{XY}}{\sigma_X \sigma_Y}.$$

- Also denoted as  $\rho_{XY}$ .
- $-1 \leq \text{Corr}(X, Y) \leq 1$
- Invariant to scale: if I double  $Y$ ,  $\text{Corr}(X, Y)$  will not change.


# Correlation Coefficient

Perfect positive correlation:  $\text{Corr}(X, Y) = 1$ .




# Correlation Coefficient

Perfect negative correlation:  $\text{Corr}(X, Y) = -1$ .




# Correlation Coefficient

Positive correlation:  $\text{Corr}(X, Y) > 0$ .




# Correlation Coefficient

Negative correlation:  $\text{Corr}(X, Y) < 0$ .



# Correlation Coefficient

No correlation:  $\text{Corr}(X, Y) = 0$ .



# Variance: Rule 03

For constants  $a$  and  $b$ ,

$$\text{Var}(aX + bY) = a^2 \text{Var}(X) + b^2 \text{Var}(Y) + 2ab \text{Cov}(X, Y).$$

- If  $X$  and  $Y$  are uncorrelated, then

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$$

- If  $X$  and  $Y$  are uncorrelated, then

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y)$$

Expanded proof

# Estimators



# Estimators

Why do we estimate things? *Because we can't measure everything*

Suppose we want to know the average height of the population in the US

- We have a sample 1 million Americans

*How can we use these data to estimate the height of the population?*

# Estimators

## **Estimand:**

Quantity that is to be estimated in a statistical analysis

## **Estimator:**

A rule (or formula) for estimating an unknown population parameter given a sample of data.

## **Estimate:**

A specific numerical value that we obtain from the sample data by applying the estimator.

# Estimators *Ex.*

Suppose we want to know the average height of the population in the US

- We have a sample 1 million Americans

**Estimand:** The population mean ( $\mu$ )


**Estimator:** The sample mean ( $\bar{X}$ )

$$\bar{X} = \frac{1}{n} \sum_{i=1}^n X_i$$

**Estimate:** The sample mean ( $\hat{\mu} = 5'6''$ )

# Properties of estimators

Imagine that we want to estimate an unknown parameter  $\mu$ , and we know the distributions of three competing estimators. *Which one should we use?*



# Properties of estimators

Question: *What properties make an estimator reliable?*

Answer (1): **Unbiasedness**

On average, does the estimator tend toward the correct value?

More formally: Does the mean of estimator's distribution equal the parameter it estimates?


$$\text{Bias}_{\mu}(\hat{\mu}) = E[\hat{\mu}] - \mu$$

# Properties of estimators


Question What properties make an estimator reliable?

AO1: **Unbiasedness**

**Unbiased estimator:**  $E[\hat{\mu}] = \mu$



**Biased estimator**  $E[\hat{\mu}] \neq \mu$



# Unbiasedness example

Is the sample mean  $\frac{1}{n} \sum_{i=1}^n x_i = \hat{\mu}$  an unbiased estimator of the population mean  $E(x_i) = \mu$ ?

$$\begin{aligned} E[\hat{\mu}] &= E\left[\frac{1}{n} \sum_{i=1}^n x_i\right] \\ &= \frac{1}{n} \sum_{i=1}^n E[x_i] \quad \} \text{ rule 3} \\ &= \frac{1}{n} \sum_{i=1}^n \mu \quad \} \text{ by definition} \\ &= \mu \end{aligned}$$

# Properties of estimators

Question What properties make an estimator reliable?

A02: **Efficiency** (*low variance*)

The central tendencies (means) of competing distributions are not the only things that matter. We also care about the variance of an estimator.

$$\text{Var}(\hat{\mu}) = E \left[ (\hat{\mu} - E[\hat{\mu}])^2 \right]$$


**Lower variance** estimators estimate closer to the mean in each sample



# Properties of estimators

Question: *What properties make an estimator reliable?*


Ans: **Efficiency** (*low variance*)



# The bias-variance tradeoff

Should we be willing to take a bit of bias to reduce the variance

In economics/causal inference we emphasize unbiasedness



# Unbiased estimators

In addition to the sample mean, there are several other unbiased estimators we will use often.

- **Sample variance** estimates variance  $\sigma^2$ .
- **Sample covariance** estimates covariance  $\sigma_{XY}$ .
- **Sample correlation** estimates the pop. correlation coefficient  $\rho_{XY}$ .

# Unbiased estimators

Sample variance,  $S_X^2$ , is an unbiased estimator of the pop. variance  $\sigma^2$

$$S_X^2 = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})^2.$$

Sample covariance,  $S_{XY}$ , is an unbiased estimator of the pop. covariance,  $\sigma_{XY}$

$$S_{XY} = \frac{1}{n-1} \sum_{i=1}^n (X_i - \bar{X})(Y_i - \bar{Y}).$$

# Unbiased estimators

Sample correlation  $r_{XY}$  is an unbiased estimator of the pop. correlation coefficient  $\rho_{XY}$

$$r_{XY} = \frac{S_{XY}}{\sqrt{S_X^2} \sqrt{S_Y^2}}.$$

# Sampling

# Sampling

## Population:

A group of items or events we would like to know about.

**Ex.** Americans, games of chess, cats in Eugene, etc.

## Parameter<sup>1</sup>

a value that describes that population

**Ex.** Mean height of American, average length of a chess game, median weight of the kitties

# Sampling

## Sample:

A survey of a subset of the population.

**Ex.** Respondents to a survey, random sample of econ students at the UO

Often we aim to draw observations randomly from the population

- Advantageous as it becomes a **representative sample** of the population...



# Sampling distributions

**Focus:** Populations vs Samples

- How can we make inferences about a **population** based on a small **sample** of the population?
- How do we learn about an unknown population parameter of interest?

**Challenge:** Usually missing data of the entire population.

**Solution:** Sample from the population and estimate the parameter.

- Draw  $n$  observations from the population, then use an estimator.

# Sampling distributions


There are myriad ways to produce a sample,<sup>1</sup> but we will restrict our attention **to simple random sampling**, where

1. Each observation is a random variable.
2. The  $n$  random variables are independent.


Life becomes much simpler for the econometrician.

# Population vs. sample

Question: *Why do we care about population vs. sample?*




Population




Population relationship

# Population vs sample

Question: *Why do we care about population vs. sample?*




10 random individuals




Population relationship

# Population vs sample

Question: *Why do we care about population vs. sample?*




10 random individuals




Population relationship

# Population vs sample

Question: *Why do we care about population vs. sample?*



10 random individuals



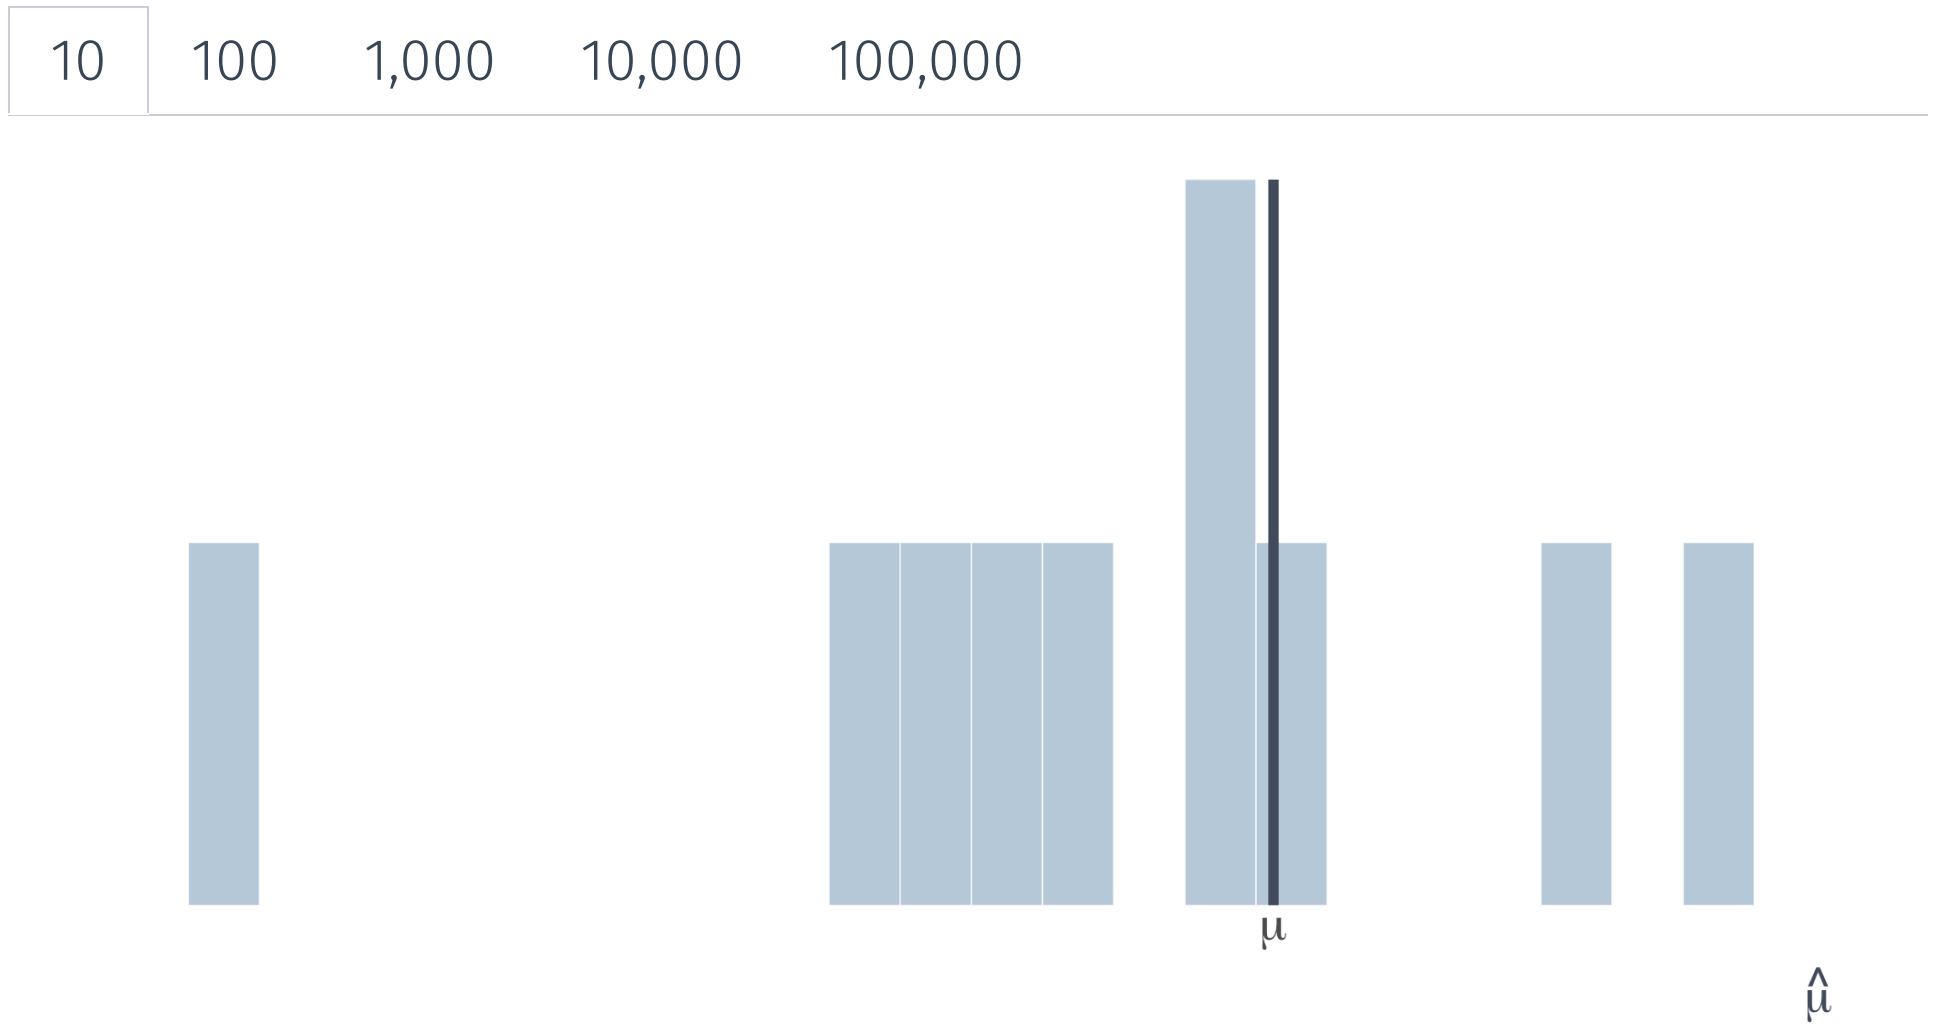
Population relationship

Let's repeat this **10,000 times** and then plot the estimates.  
(This exercise is called a Monte Carlo simulation.)

How in the world do I do that

► Show the code






Regular resampling means of 10 obs at a time

# Population vs. sample

Question: *Why do we care about population vs. sample?*



As the number of samples approach infinity

On average, the mean of the samples are close to the population mean

- Some individual samples can miss the mark.
- The difference between individual samples and the population creates **uncertainty**

# Population vs. sample

Question: *Why do we care about population vs. sample?*

Answer: Uncertainty matters.

- $\hat{\mu}$  is a random variable that depends on the sample.
- We don't know if our sample is representative of the population.
- Individual sample means can be biased
- We have to keep track of this uncertainty.

# Population distributions

**Consider the following argument** (this slide scrolls down)

Suppose we have some estimator  $\hat{\theta}$  for a parameter  $\theta$ :

- $\theta$  is unobserved, but assume  $\hat{\theta}$  follows a probability distribution  $p(\hat{\theta})$
- We hypothesize some value, say  $\theta = 2.5$
- We use our estimator  $\hat{\theta}$  to calculate an estimate.  $\hat{\theta} = 45$
- If we make an **assumption** of the distribution of  $\hat{\theta}$ , we can calculate the probability of getting  $\hat{\theta} = 45$  when  $\theta = 2.5$  is true.
- For sake of argument, let's say that the probability that  $\theta = 2.5$  if we observe  $\hat{\theta} = 45$  is less than 0.001

We can say

if  $\theta$  really was 2.5, then the probability of getting  $\hat{\theta} = 45$  is super super low. Thus the probability that  $\theta$  is actually 2.5 is super super low”.

- We can make statements about the true value of  $\theta$  just by knowing the distribution of our preferred estimator  $\hat{\theta}$

But what distribution should we be assuming?

# The Central Limit Theorem

## Theorem

Let  $x_1, x_2, \dots, x_n$  be a random sample from a population with mean  $E[X] = \mu$  and variance  $\text{Var}(X) = \sigma^2 < \infty$ , let  $\bar{X}$  be the sample mean. Then, as  $n \rightarrow \infty$ , the function  $\frac{\sqrt{n}(\bar{X} - \mu)}{S_x}$  converges to a Normal Distribution with mean 0 and variance 1.

- CLT states that when  $n \rightarrow \infty$ , the sample mean will be normally distributed.
- The Law of Large Number (LLN) states that as  $n \rightarrow \infty$ , the sample converges on the population mean.

# The Central Limit Theorem

## **Some interesting YouTube links:**

- A more in depth explanation + visualization
- What is so special about the normal distribution?

# Data types



# Data

There are **two** broad types of data

## 1. **Experimental data**

Data generated in controlled, laboratory settings<sup>1</sup>

Ideal for **causal identification**, but difficult to obtain

- Logistically intractable
- Expensive
- Morally repugnant

# Data

There are **two** broad types of data

1. **Experimental data**
2. **Observational data**

Data generated in non-experimental settings

Types of observational data:

- Surveys
- Census
- Administrative data
- Environmental data
- Transaction data
- Text and image data

Commonly used though poses challenges to **causal identification**

# Data types: Cross sectional

Sample of individuals from a population at a point in time

Ideally collected using **random sampling**

- **random sampling + sufficient sample size = representative sample**
- Non-random sampling is more common and difficult to work with

Note: Used extensively in applied microeconomics<sup>1</sup> and is the main focus of this course

1. Applied microeconomics = Labor, health, education, public finance, development, industrial organization and urban economics

# Data types: Time series

Observations of variables over time

- Ex.**
- Quarterly GDP
  - Annual infant mortality rates
  - Daily stock prices

Complication: Observations are not independent draws

- eg GDP this quarter is highly correlated to GDP last quarter

More advanced methods needed<sup>1</sup>

<sup>1</sup> See EC 421 and EC 422

# Data types: Pooled cross sectional

Cross sections from different points in time

Useful for studying relationship that change over time.

Again, requires more advanced methods<sup>1</sup>

# Data types: Panel data

Time series for each cross sectional unit

**Ex.** Daily attendance across my class

Can control for unobserved characteristics

Again, requires more advanced methods<sup>1</sup>

# Data types: Messy data

Analysis ready dataset are rare. Most data are *messy*

**Data wrangling** is a non-trivial part of an economist or data scientist/analyst's job

 has a suite of packages that facilitate data wrangling:

- The `tidyverse`: `readr`, `tidyr`, `dplyr`, `ggplot2` + others

# Table of Contents

## Admin

1. Admin

## Review

1. Notation
2. Basic probability
3. Distributions
4. Moments
5. Estimators
6. Sampling
7. Data types



# Appendix

# Variance: Rule 03 Expanded

## Back to Variance Rule 03

The variance of a random variable  $X$  is defined as:

$$\text{Var}(X) = E[(X - \mu_X)^2]$$

$\text{Cov}(X, Y)$  is defined as:

$$\text{Cov}(X, Y) = E[(X - \mu_X)(Y - \mu_Y)]$$

For two random variables  $X$  and  $Y$ , the variance of their sum  $X + Y$  is:

$$\text{Var}(X + Y) = E[((X + Y) - (\mu_X + \mu_Y))^2]$$

Expanding the squared term, we get:

$$\begin{aligned}
\text{Var}(X + Y) &= E[(X - \mu_X + Y - \mu_Y)^2] \\
&= E[(X - \mu_X)^2 + 2(X - \mu_X)(Y - \mu_Y) + (Y - \mu_Y)^2] \\
&= E[(X - \mu_X)^2] + E[2(X - \mu_X)(Y - \mu_Y)] + E[(Y - \mu_Y)^2] \\
&= \text{Var}(X) + 2\text{Cov}(X, Y) + \text{Var}(Y)
\end{aligned}$$

If  $X$  and  $Y$  are uncorrelated, then  $\text{Cov}(X, Y) = 0$ , and the above simplifies to:

$$\text{Var}(X + Y) = \text{Var}(X) + \text{Var}(Y)$$

Similarly, the variance of the difference  $X - Y$  is:

$$\text{Var}(X - Y) = E[((X - Y) - (\mu_X - \mu_Y))^2]$$

Expanding the squared term, just like before:

$$\begin{aligned}
\text{Var}(X - Y) &= E[(X - \mu_X - (Y - \mu_Y))^2] \\
&= E[(X - \mu_X)^2 - 2(X - \mu_X)(Y - \mu_Y) + (Y - \mu_Y)^2] \\
&= \text{Var}(X) - 2\text{Cov}(X, Y) + \text{Var}(Y)
\end{aligned}$$

Again, if  $X$  and  $Y$  are uncorrelated,  $\text{Cov}(X, Y) = 0$ , and we have:

$$\text{Var}(X - Y) = \text{Var}(X) + \text{Var}(Y)$$